

# General Concepts

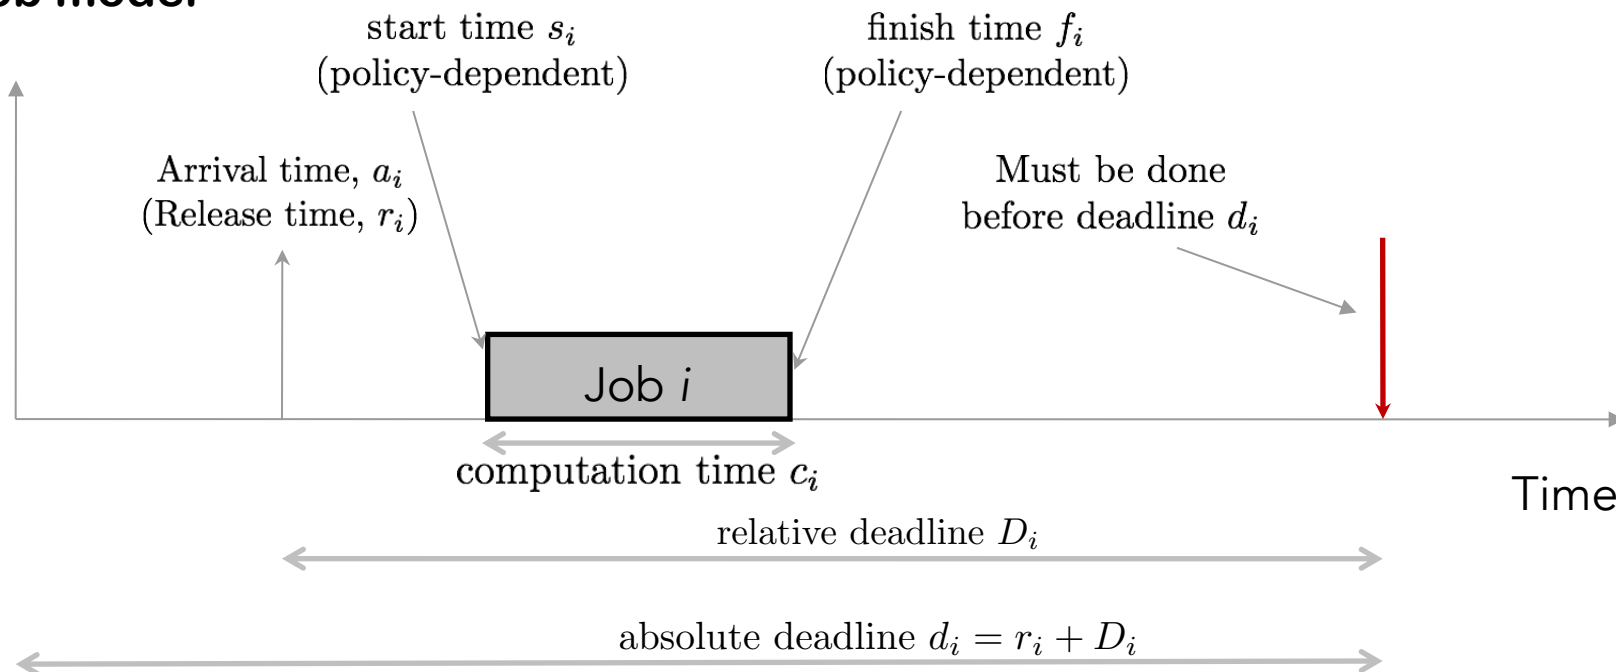
---

Introduction to real-time systems

# Abstraction

- Need an abstraction to reason about timeliness that is:
  1. Expressive enough to include realistic workload models and parameters
  2. Constrained enough to derive useful results (tractable)

## Job model



# Online vs. Offline scheduling

---

- What do we know about the jobs to be scheduled?
- **Offline**: All job parameters  $r_i, c_i, D_i, \dots$  and number of jobs  $n$  are known a priori
- **Online**: Jobs arrive as a *stream*
  - Might not know arrival times
  - We do not know number of jobs
  - We might know execution time either when job arrives or even after job finishes execution!
- In this course, mostly concerned with **offline** scheduling problems

# Job Scheduling: Policies

---

- A rule that specifies which job occupies which processor at every time instant
- May be
  - Deterministic
  - Randomized
  - Preemptive or non-preemptive
  - *Priority-driven* or *clock-driven* (or a hybrid of both, or none of which!)

# Job Scheduling: Policies

---

- Deterministic scheduling policy:  $n$  jobs

A function  $S: \mathbb{R}_+ \rightarrow \{1, \dots, n\}$  where  $S(t)$  is the index of the job that occupies the processor at time  $t$

- $S(t) = 0$  if the processor is kept idle
- Some authors impose the additional requirement that a policy be a *step function*
  - $\forall t > 0, \exists t_1, t_2 \geq 0$  with  $t \in [t_1, t_2)$  such that  $S(t') = S(t)$  for every  $t' \in [t_1, t_2)$

# Job Scheduling: Priority-driven Policies

---

- Priority-driven policies
  - Assign priorities to jobs so that the ready job with the highest priority occupies the processor
  - Priority can be static or vary through time as job execution advances (dynamic priority)
  - **Q:** With  $n$  jobs, how may **static priority assignments exist** where each job has a unique priority?
  - **Q:** Can an optimal static priority assignment be found in time that is polynomial in the problem parameters?
  - What does **optimal** mean?
    - For now, minimizes some scheduling objective → will consider more sophisticated definitions of optimality shortly
- Also called **work-conserving, list scheduling, greedy** → processor is never left idle on purpose, always schedule a job if one is available

# Job Scheduling: Randomized Policies

---

- Essentially a stochastic process  $\pi = \{\pi_t: t\}$
- At  $t > 0$ ,  $\pi_t$  is a conditional probability distribution on jobs given past data
- Past data includes previous scheduling decisions, job allocations so far, etc
- Intuitively, the policy throws a coin to aid in making allocation decisions
- Mostly used in stochastic job models where some job parameters are themselves random variables (later)

# Job Scheduling: Objectives (cost functions)

We have  $n$  jobs  $J_1, \dots, J_n$  with parameters  $r_i, c_i, D_i$ , etc

- **Objective related to timing properties**

- minimize *total completion (finish) time*:  $\sum_{i=1}^n f_i$
- minimize total **weighted** completion time:  $\sum_{i=1}^n w_i f_i$
- Minimize **makespan**  $\max_i f_i$ 
  - *Makespan*: finish time of last job to leave the system (length of schedule)
  - Is “machine owner oriented”: A minimum makespan implies a good utilization of the machine(s).
  - Not suitable for interactive applications: Jobs released early might be delayed till end of an optimal makespan schedule → not acceptable by “users”
- Minimize total flow (*response*) time:  $\sum_{i=1}^n (f_i - r_i)$ 
  - $(f_i - r_i)$ : total time job  $i$  is in system = *waiting time* + *processing time*
  - Beneficial from “users” perspective
  - Also minimizes total completion time and total waiting time
  - **Drawback**: Can lead to *starvation*: Might still cause some jobs to be delayed unboundedly



# Job Scheduling: Objectives (cost functions)

- Minimize max. *flow time*:  $\max_i (f_i - r_i)$ 
  - Schedule responsive to each job
  - Starvation-free
- Minimize *average flow (response) time*:  $\frac{1}{n} \sum_{i=1}^n (f_i - r_i)$ 
  - Used in *soft* real time systems to improve QoS
- **Lateness**:  $L_i = f_i - d_i \rightarrow$  Minimize maximum lateness  $L_{\max} = \max_i L_i$ 
  - Worst case violation of deadlines. The earlier the better!
- **Tardiness**:  $T_i = \max\{0, L_i\} \rightarrow$  Minimize maximum tardiness  $T_{\max} = \max_i T_i$ 
  - Doesn't care how early (before the deadline) jobs finish
- **# of tardy jobs**  $\sum_{i=1}^n \mathbf{1}\{L_i > 0\} = \sum_{i=1}^n \mathbf{1}\{f_i > d_i\}$

**Deadline-related**

# Job Scheduling: Stochastic Models

---

- What if we have only **statistical data** of job parameters? E.g.,
  - Jobs (events) *arrive* randomly
  - Job execution time (**demand**) is random
    - **Inventory** where product **demands** are not known exactly but we know their distribution
    - Machine might *fail* randomly, causing execution times to be “stretched” randomly by (down time + how long it takes to service machine)
- **Questions of interest:**
  - What is the *probability* that jobs miss their deadlines under some policy?
  - What is the *expected* max lateness  $\mathbb{E}L_{\max}$  under some fixed policy?
  - What policy *minimizes* the expected max lateness?

# Job Scheduling: Preemption

---

- Can we interrupt a job once it starts execution?
- Preemptible jobs/resources
  - **Example:** execution of programs on your computer may be interleaved to implement multitasking
- Non-preemptible jobs/resources:
  - You might be on an important phone call so you block all other incoming calls
  - Disable interrupts during the execution of an ISR
  - When using resources with complex or costly state associated with it → prohibitive context switching costs → *blocks on disk*
- No-preemption problems are usually computationally harder (**NP**-Hard)
- Allowing preemption gives more flexibility to scheduler → sometimes more tractable
  - **Example:** Minimizing maximum lateness with arbitrary release times is **NP**-hard when *no* preemptions allowed → becomes **poly-time** solvable when preemptions allowed
  - **Downside:** preemption (context switching) costs might be prohibitive

# Example: Minimize total weighted completion time

---

- Given have  $n$  jobs  $J_1, \dots, J_n$ 
  - $J_i$  has execution time  $e_i > 0$  and weight (importance)  $w_i \geq 0$
  - All jobs arrive at time 0  $\rightarrow r_i = 0 \quad \forall i \in \{1, \dots, n\}$
  - One processor
- **For instance:** Job might be *message to be broadcasted*, all messages ready for transmission at time 0, message  $i$  has length  $e_i > 0$ , and a fraction  $w_i \geq 0$  of receivers are interested in message  $i$ 
  - Finish time  $f_i$  is the time at which message is fully transmitted

$$\text{minimize } \sum_{i=1}^n w_i f_i$$

# Example: Minimize total weighted completion time


---

- Is there an optimal priority assignment?
- If so, how does the priority function look like?
- Think ... greedy
- Key: Job interchange argument
- Will do proof of optimality in class

## Example: $1|\text{prmt}, r_j|L_{\max}$

**Problem:** Schedule  $n$  jobs with *arbitrary arrival times*  $r_1, \dots, r_n$  and *relative deadlines*  $D_1, \dots, D_n$  on a single processor with preemption allowed to minimize **maximum lateness**

- $1|\text{prmt}, r_j|L_{\max}$  is an example of the **Graham notation** for job scheduling problems
- **Earliest Deadline First (EDF)** is optimal  $\rightarrow$  At every time instant, schedule the job whose *absolute deadline*  $d_j = r_j + D_j$  is nearest if job  $j$  has already arrived ( $t \geq r_j$ )
- **Preemptive:** if job  $i$  is running when job  $j$  arrives at time  $r_j$  and  $d_j < d_i$ , then remove job  $i$  from processor and start running job  $j$ , otherwise keep running job  $i$



# Optimality of preemptive EDF for $1|prmt, r_j|L_{\max}$

---

**Problem:** Schedule  $n$  jobs with *arbitrary arrival times*  $r_1, \dots, r_n$  and relative deadlines  $D_1, \dots, D_n$  on a single processor with preemption allowed to minimize **maximum lateness**

- Optimality of preemptive EDF follows by the two following claims:
- **Claim1:**  $L_{\max} \geq r(S) + c(S) - d(S)$  for any  $S \subset \{1, \dots, n\}$ , where  
 $r(S) = \min_{j \in S} r_j$ ,  $c(S) = \sum_{j \in S} c_j$ ,  $d(S) = \max_{j \in S} d_j$  (regardless of scheduling algorithm)
- **Claim2:** preemptive EDF gives a schedule with  $L_{\max}(\text{EDF}) = \max_{S \subset \{1, \dots, n\}} r(S) + c(S) - d(S)$   
(achieves the min possible  $L_{\max}$ )

# Notes on preemptive EDF

- Is an *online* policy
- How many preemptions for  $n$  jobs?
  - equals the number of **distinct** *release times*
- Can be implemented in  $O(n \log n)$



## Stochastic $1||L_{\max}$

- Suppose job execution times are independent positive random variables on some probability space
- Which rule minimizes expected max lateness  $\mathbb{E}L_{\max}$  ?
- Hint: very similar to deterministic case!


# Schedulability problem for non-preemptive jobs with arbitrary release times

- We have  $n$  **non-preemptible** jobs  $J_1, \dots, J_n \rightarrow J_i$  has parameters  $(c_i, r_i, D_i)$
- **Problem:** *how to schedule jobs **non-preemptively** so that all jobs meet their deadlines if this is possible.*
- *There is no cost function!*
  - *Such problems are called **feasibility** (or **schedulability**) problems*
- How do we define optimality then?
  - A rule  $S$  is **optimal** in the following sense: ***For every instance of the problem, it can always meet the deadlines if any other rule can***
  - *Equivalent: If rule  $S$  cannot meet deadlines, then no other rule can*
- Is non-preemptive EDF optimal?

# Non-preemptive scheduling with arbitrary release times

We have  $n$  **non-preemptible** jobs  $J_1, \dots, J_n \rightarrow J_i$  has parameters  $c_i, r_i, D_i$


- **Problem:** how to schedule jobs *non-preemptively* so that all jobs meet their deadlines if this is possible.
- Is non-preemptive EDF optimal?



# Non-preemptive scheduling with arbitrary release times

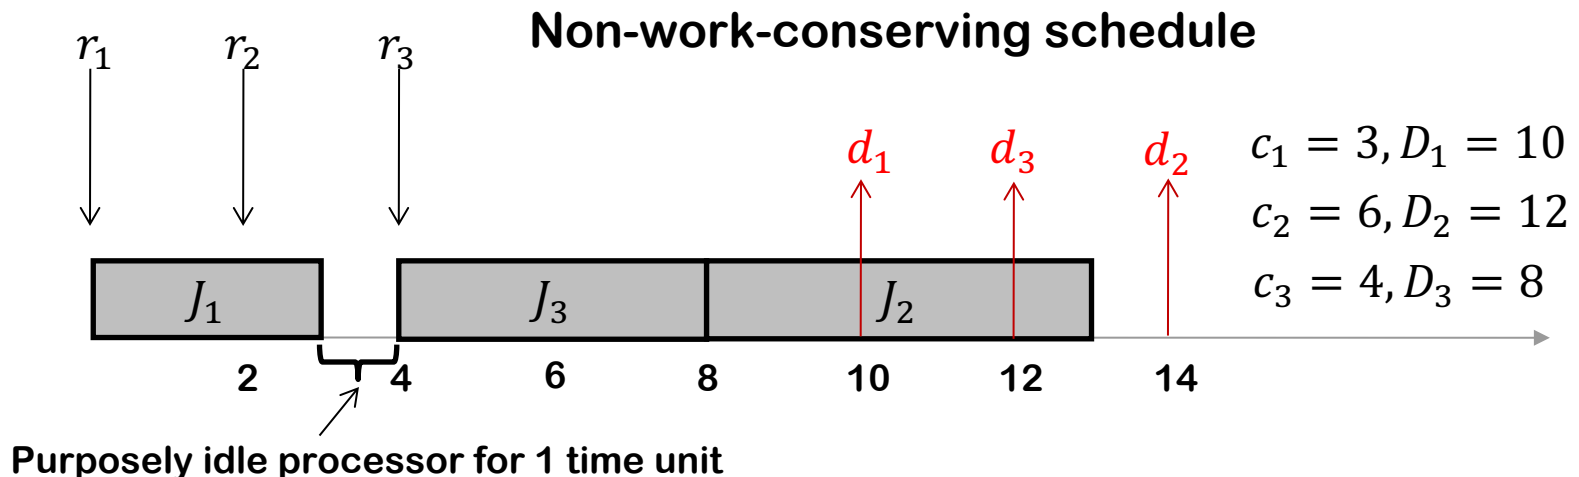
We have  $n$  **non-preemptible** jobs  $J_1, \dots, J_n \rightarrow J_i$  has parameters  $c_i, r_i, D_i$

- **Problem:** *how to schedule jobs **non-preemptively** so that all jobs meet their deadlines if this is possible.*
- Is there a policy under which this job set can meet its deadlines?



# Non-preemptive scheduling with arbitrary release times


- This schedule introduces idle time on purpose to meet deadlines
  - Schedule is *non-work-conserving*
- **Synonyms:** work-conserving, list schedule, greedy, priority-driven: Algorithms that *never leave the processor idle*
  - They dispatch a job to the processor as soon as one is available!
- **This example shows more:** *no priority-driven algorithm is optimal for the non-preemptive problem with arbitrary deadlines, execution times, and release times*



# A scheduling anomaly


We have  $n$  **non-preemptible** jobs  $J_1, \dots, J_n \rightarrow J_i$  has parameters  $c_i, r_i, D_i$

- **Problem:** how to schedule jobs *non-preemptively* so that all jobs meet their deadlines if this is possible.
- What happens under EDF if we increase  $c_1$  from 3 to 4?
  - *Intuition:* More workload  $\rightarrow$  more deadlines should be missed under same policy
  - *Reality:* All jobs meet their deadlines!
  - *Question:* Is this policy *predictable*?



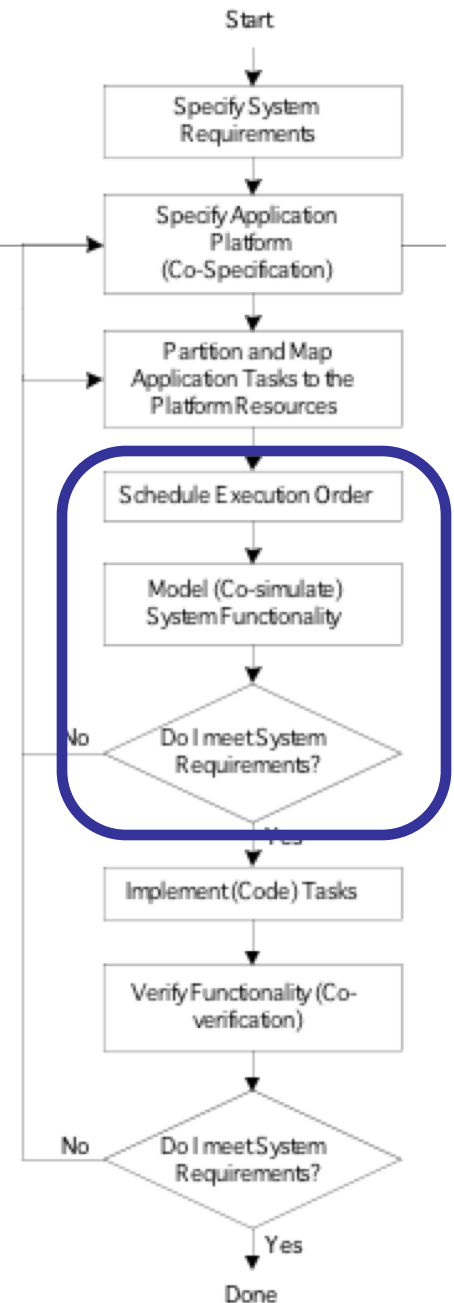
# Review

- What is a real-time system?
- What is an embedded system?
- What characteristic of a real-time system is probably the most important?



# The system design process


- Designing any computer system involves many steps.
- Some steps are common to many types of systems.
- A few steps are more important in a real-time system.
  - **Scheduling** is one such operation.
  - How do we know if a set of tasks can be scheduled in a predictable manner?
- We will touch upon other parts of the design process later in the course.





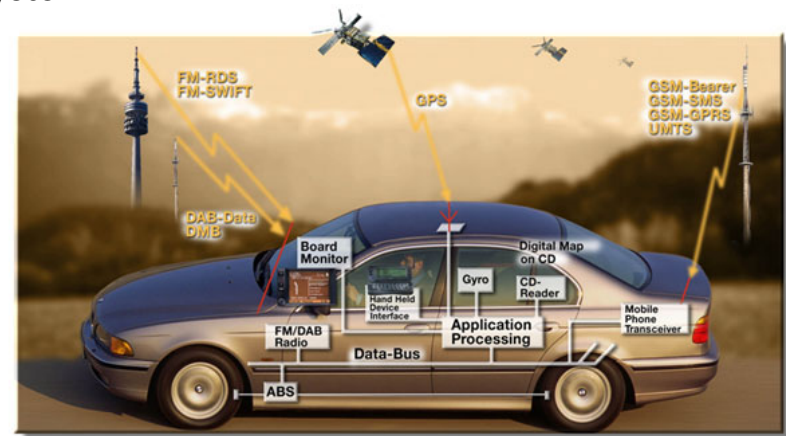
# The schedulability question: Drive-by-Wire Example

- Consider a control system in a future vehicle
  - Steering wheel sampled every 10 ms – wheel positions adjusted accordingly (computing the adjustment takes 4.5 ms of CPU time)
  - Brakes sampled every 4 ms – break pads adjusted accordingly (computing the adjustment takes 2ms of CPU time)
  - Velocity is sampled every 15 ms – acceleration is adjusted accordingly (computing the adjustment takes 0.45 ms)
  - For safe operation, adjustments must always be computed before the next sample is taken
- Is it possible to always compute all adjustments in time?
- The underlying computer system is a **uniprocessor** system.




# The schedulability question: Drive-by-Wire Example

- Consider a control system in a future vehicle
  - Steering wheel sampled every 10 ms – wheel positions adjusted accordingly (computing the adjustment takes 4.5 ms of CPU time)
  - Brakes sampled every 4 ms – break pads adjusted accordingly (computing the adjustment takes 2ms of CPU time)
  - Velocity is sampled every 15 ms – acceleration is adjusted accordingly (computing the adjustment takes 0.45 ms)
  - For safe operation, adjustments must always be computed before the next sample is taken
- Is it possible to always compute all adjustments in time?
- The underlying computer system is a **uniprocessor** system.



# Tasks

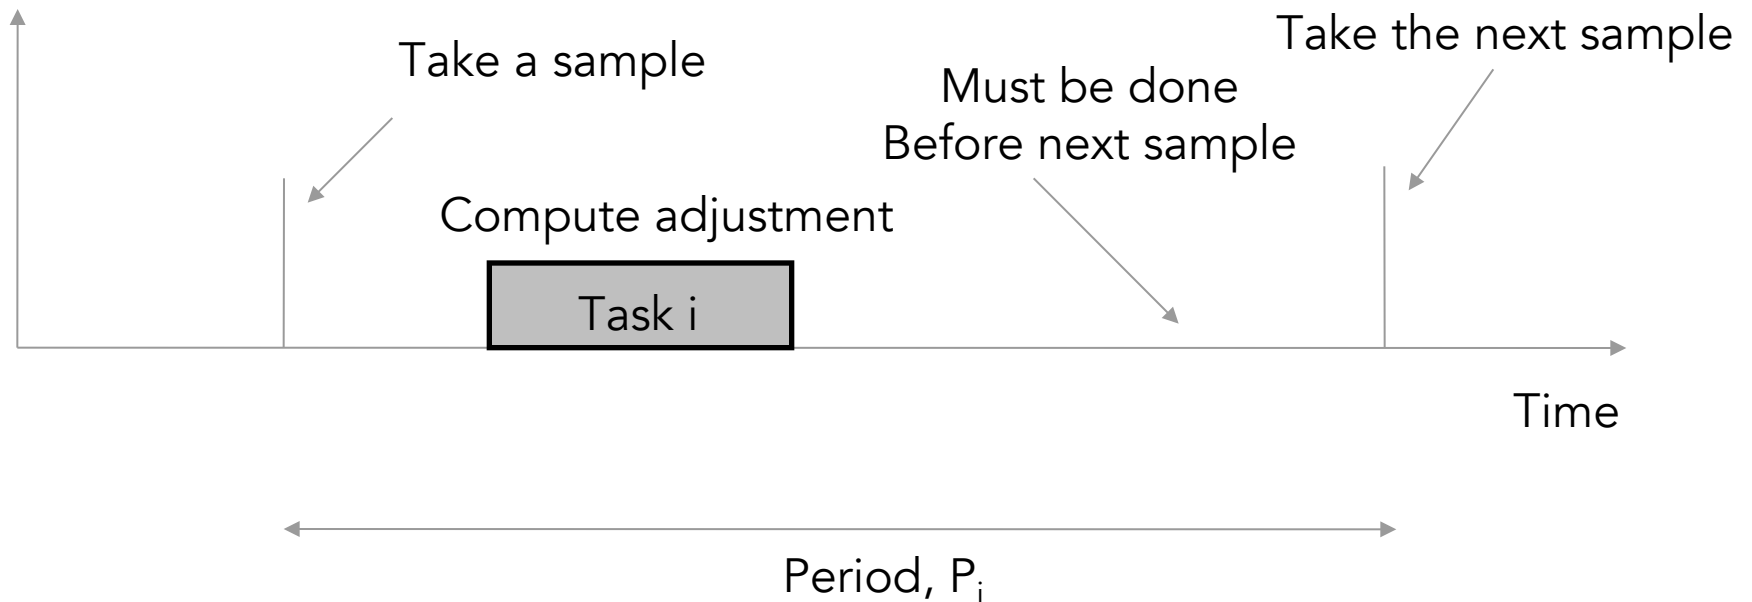
- Basic RT model: Periodic **Task model**
  - A **task** releases an infinite (indefinitely long) succession of **jobs**



# Some terminology

---

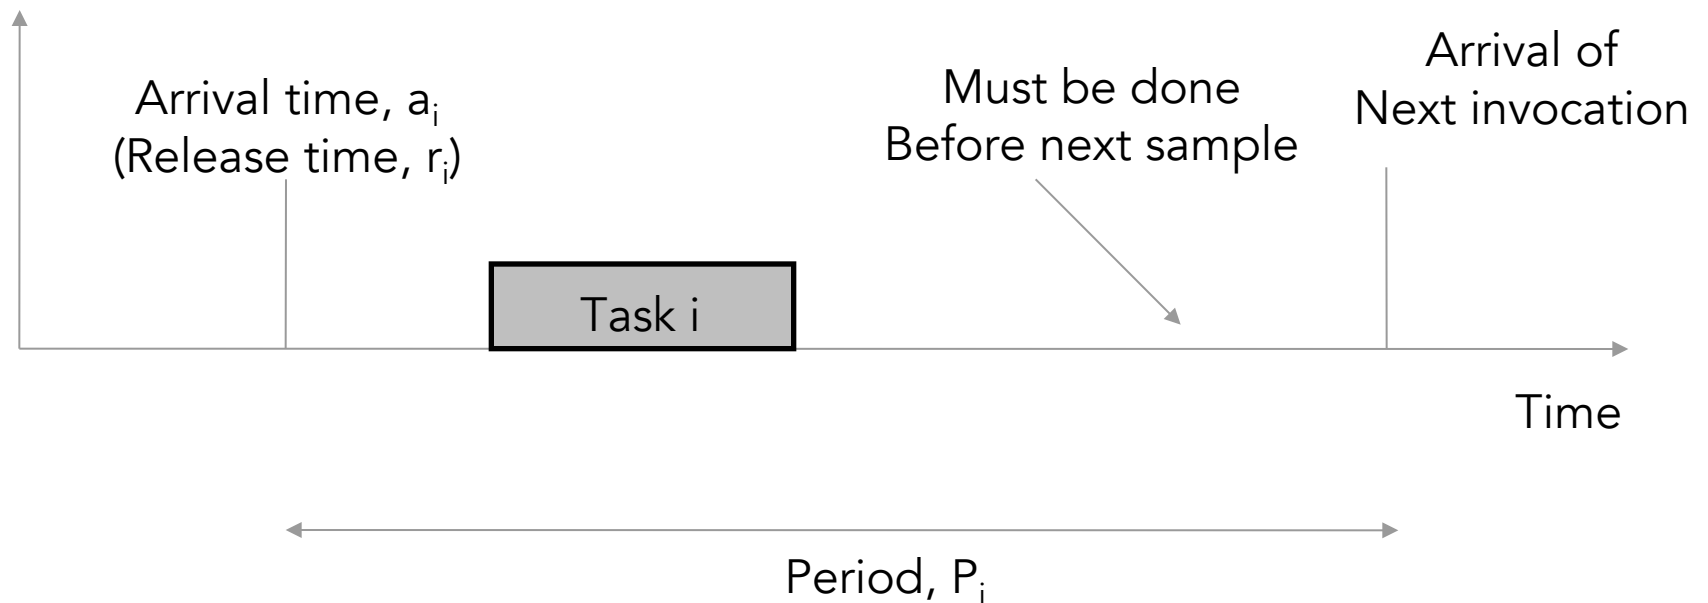
- Tasks, periods, arrival-time, deadline, execution time, etc.



# Some terminology

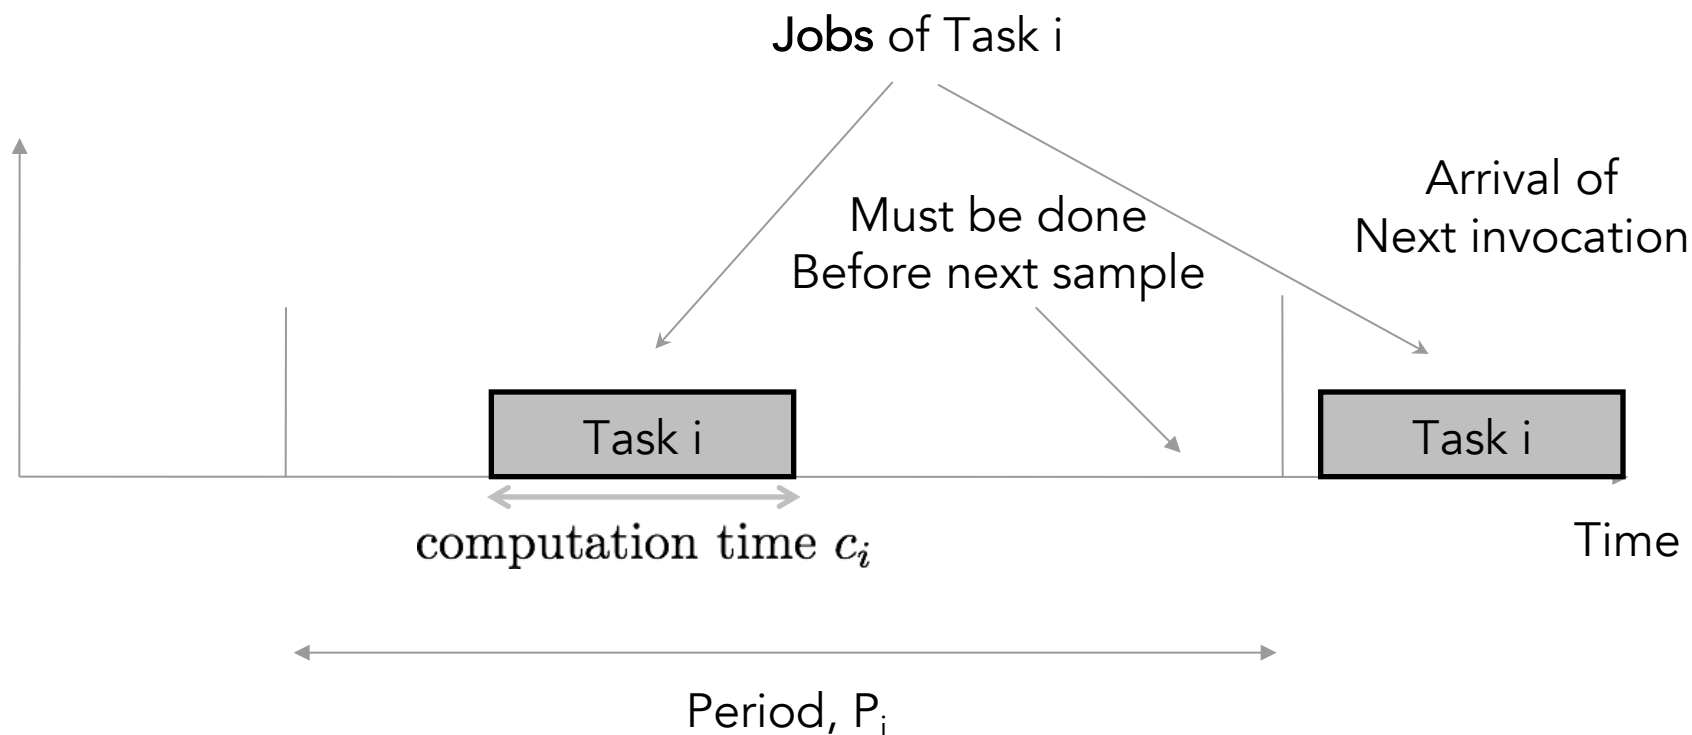
---

- Tasks, periods, arrival-time, deadline, execution time, etc.



# Some terminology

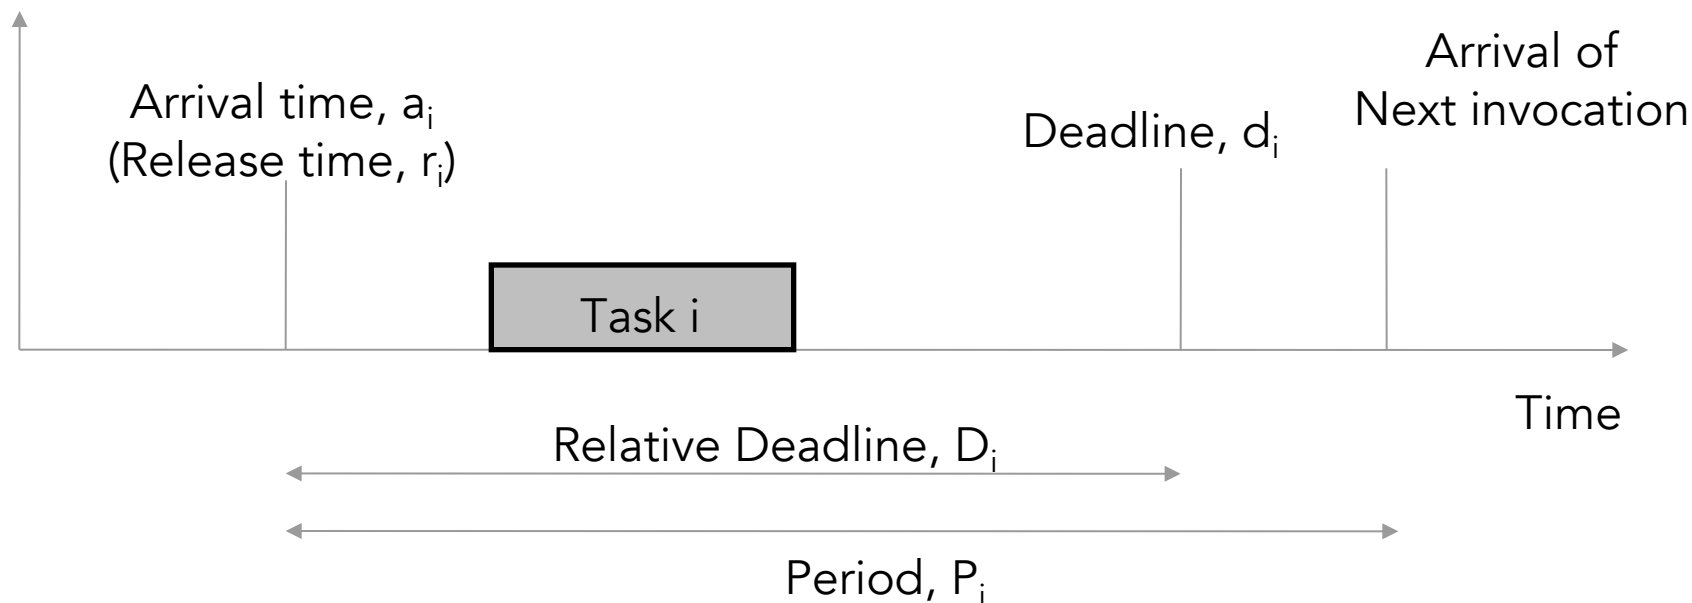
- Tasks, periods, arrival-time, deadline, execution time, etc.
- Each invocation of a task is called a “job.”
- A common assumption is that arrival times for the first job of all tasks is 0.



# Some terminology

---

- Tasks, periods, arrival-time, deadline, execution time, etc.

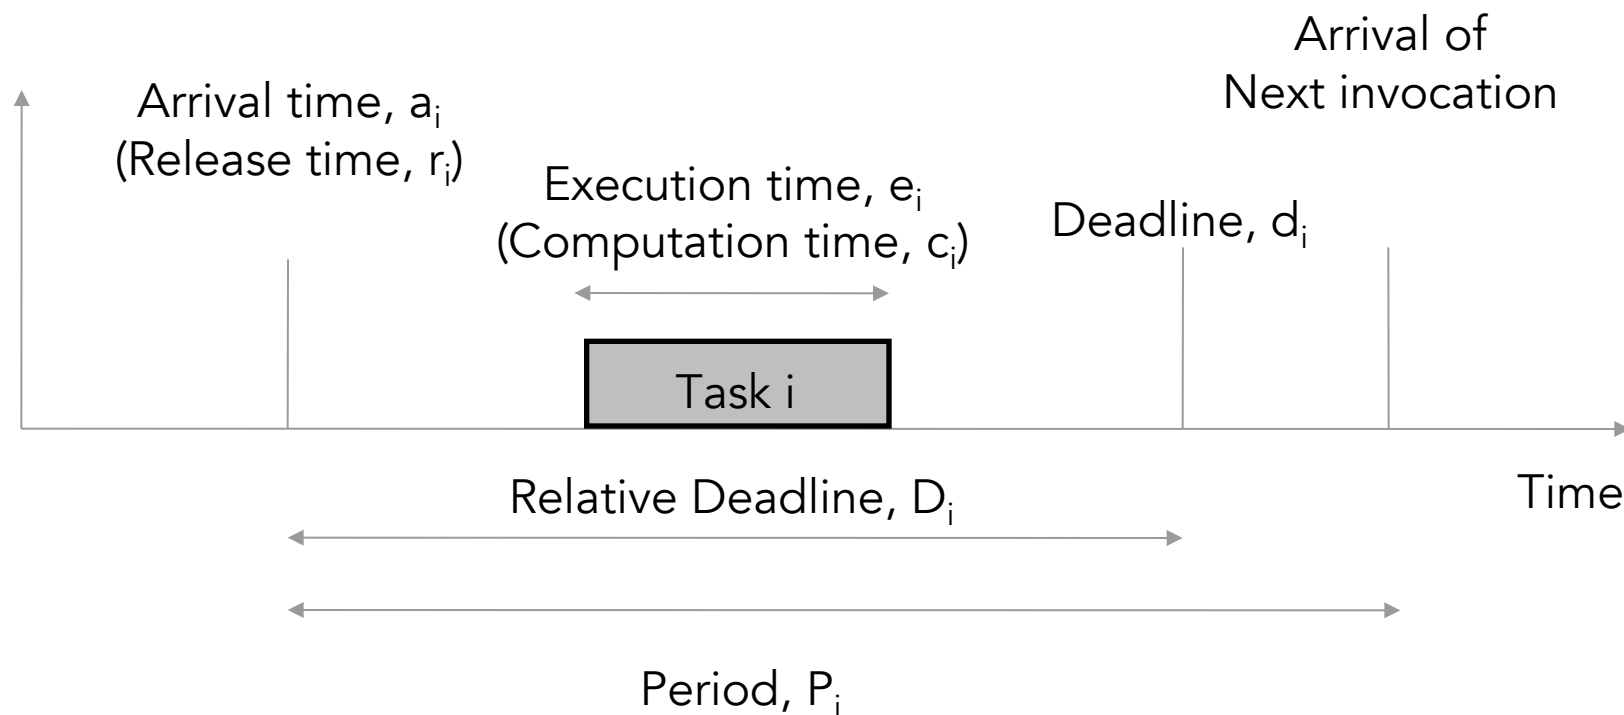


(absolute deadline)  $d_i = (\text{release time}) r_i + (\text{relative deadline}) D_i$

# Some terminology

---

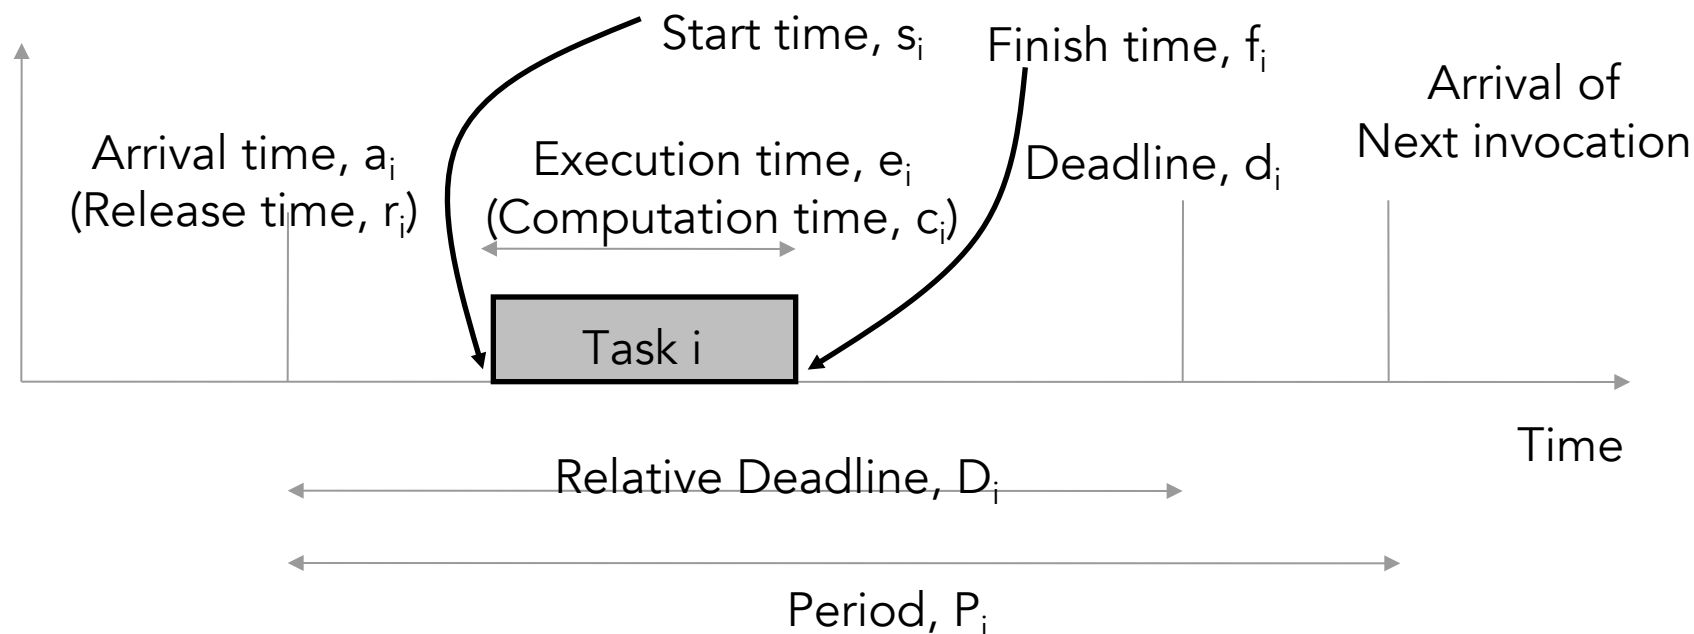
- Tasks, periods, arrival-time, deadline, execution time, etc.





# Some terminology

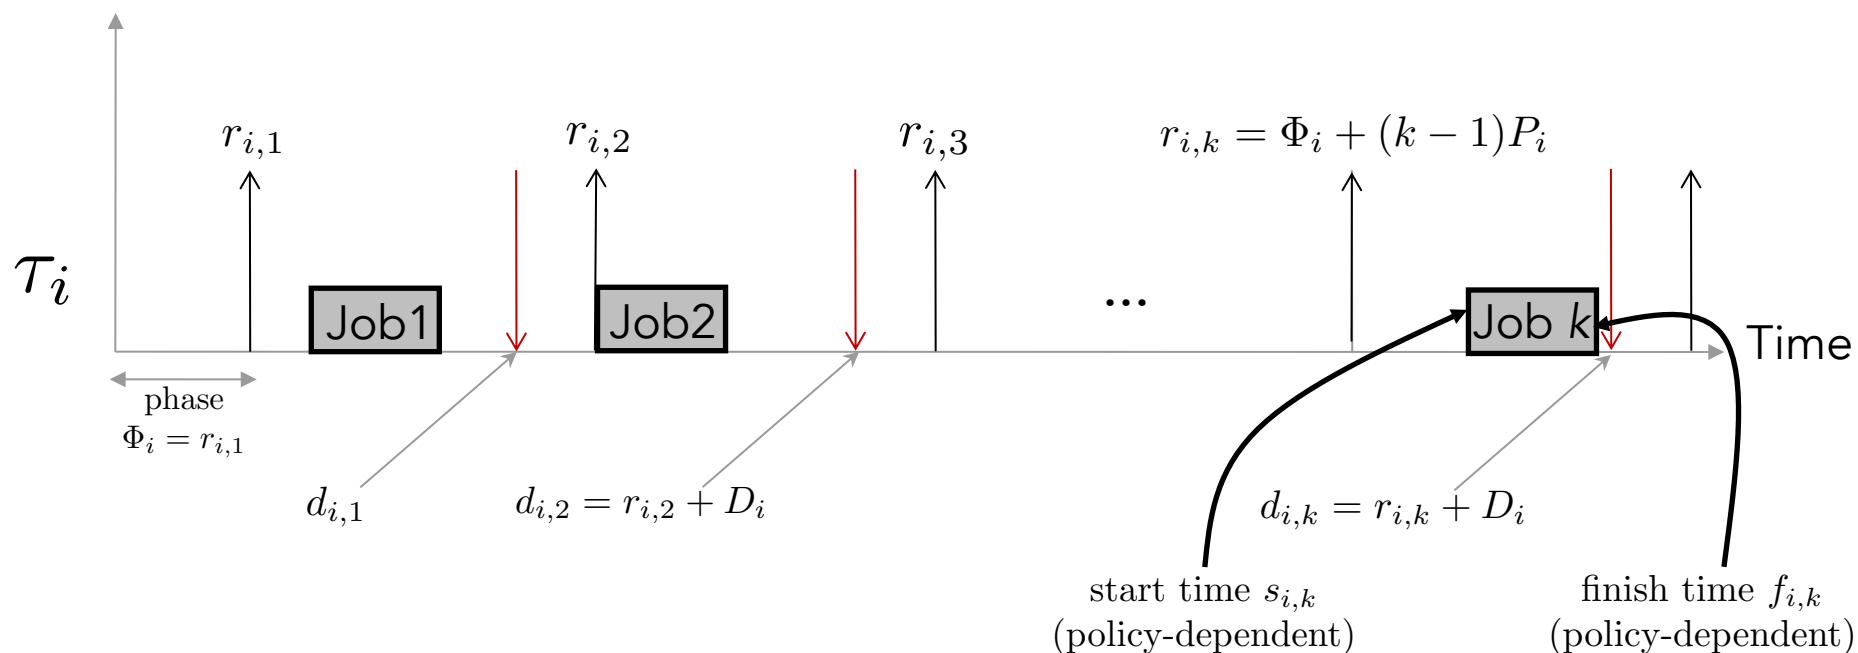
- Tasks, periods, arrival-time, deadline, execution time, etc.



# Some terminology

Tasks, periods, arrival-time, deadline, execution time, etc.

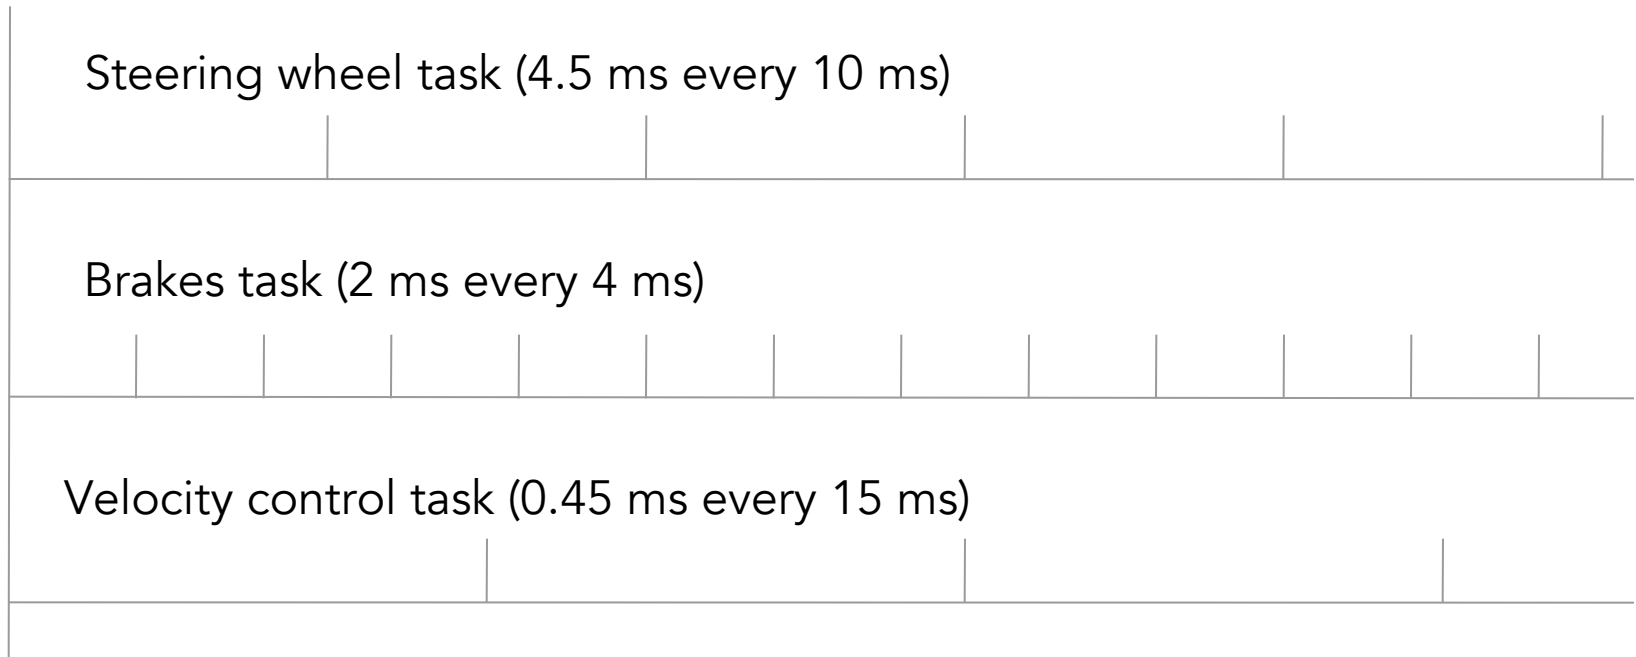
Each invocation of a task is called a "job."



# Back to the Drive-by-Wire example

---

- Find a schedule that makes sure all task invocations meet their deadlines
- Often, **relative deadlines are equal to the period lengths**



# Back to the Drive-by-Wire example

---

- Sanity check #1: Is the processor over-utilized? (e.g., if you have 5 assignments due this time tomorrow and each takes 6 hours, then  $5 \times 6 = 30 > 24 \rightarrow$  you are overutilized)
  - Hint: Check if processor utilization  $> 100\%$

Steering wheel task (4.5 ms every 10 ms)

Brakes task (2 ms every 4 ms)

Velocity control task (0.45 ms every 15 ms)

# Utilization of a task set

---


- For a set of tasks  $\{T_i\}$  with execution times  $\{e_i\}$  and periods  $\{P_i\}$ , the utilization,  $U$ , is the fraction of time, in the long run, for which the task set will use the system.

$$U = \sum_i \frac{e_i}{P_i}$$

# Task scheduling

---


- In what order should tasks be executed?
  - Hand-crafted schedule (fill timeline by hand)
  - Cyclic executive scheduling



# Task scheduling

---

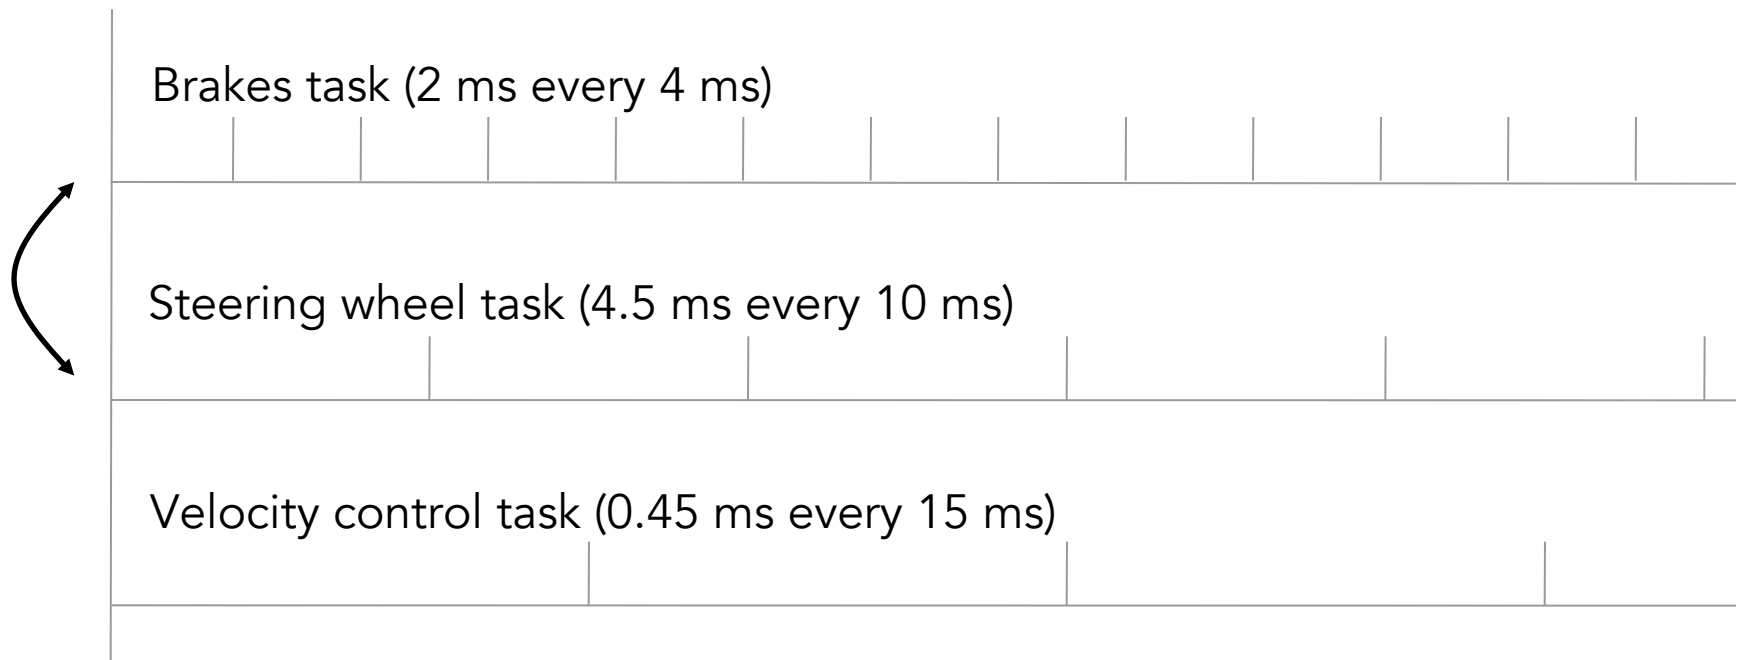
- Cyclic executive scheduling
  - Why is it called a “cyclic” executive?
  - What are the problems with cyclic executive scheduling?
    - Hard to adjust the schedule if tasks change
    - Difficult to specify



# Task scheduling

---

- In what order should tasks be executed?
  - Cyclic executive scheduling or
  - Priority based schedule (assign priorities; schedule is implied)



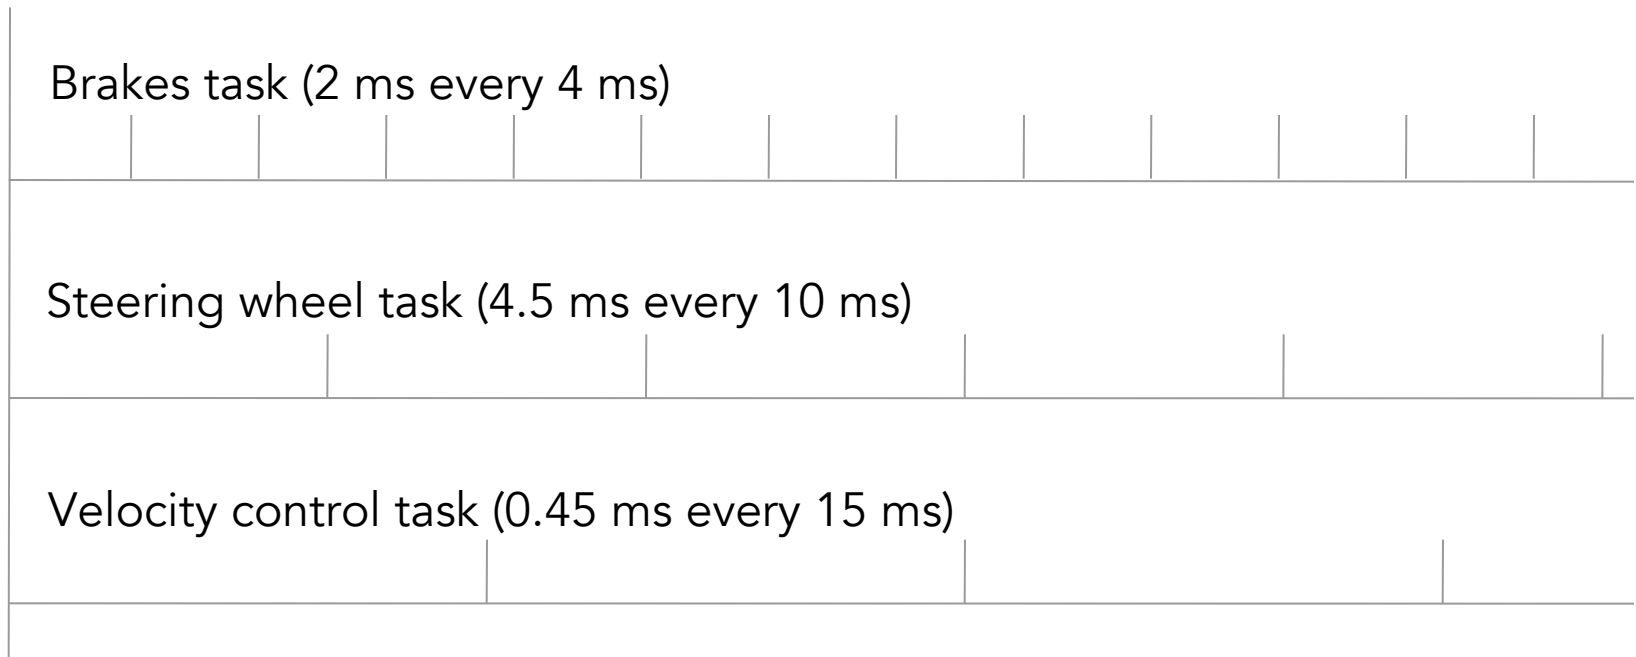
**Intuition: Urgent tasks should be higher in priority**



## Task scheduling: **Preemptive** versus **non-preemptive**?

---

- Preemptive: Higher-priority tasks can interrupt lower-priority ones
- Non-preemptive: They can't

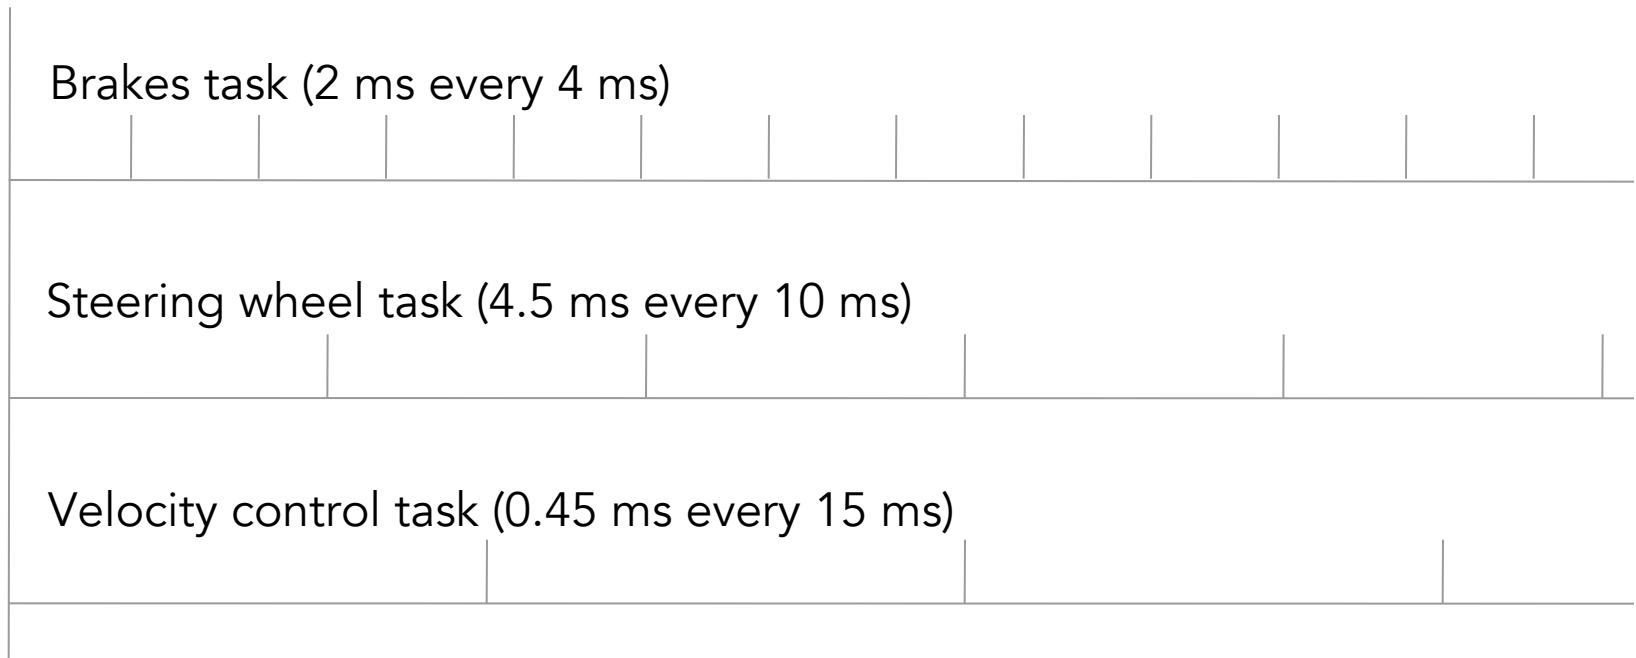


**In this example, will non-preemptive scheduling work?**

# Task scheduling

---

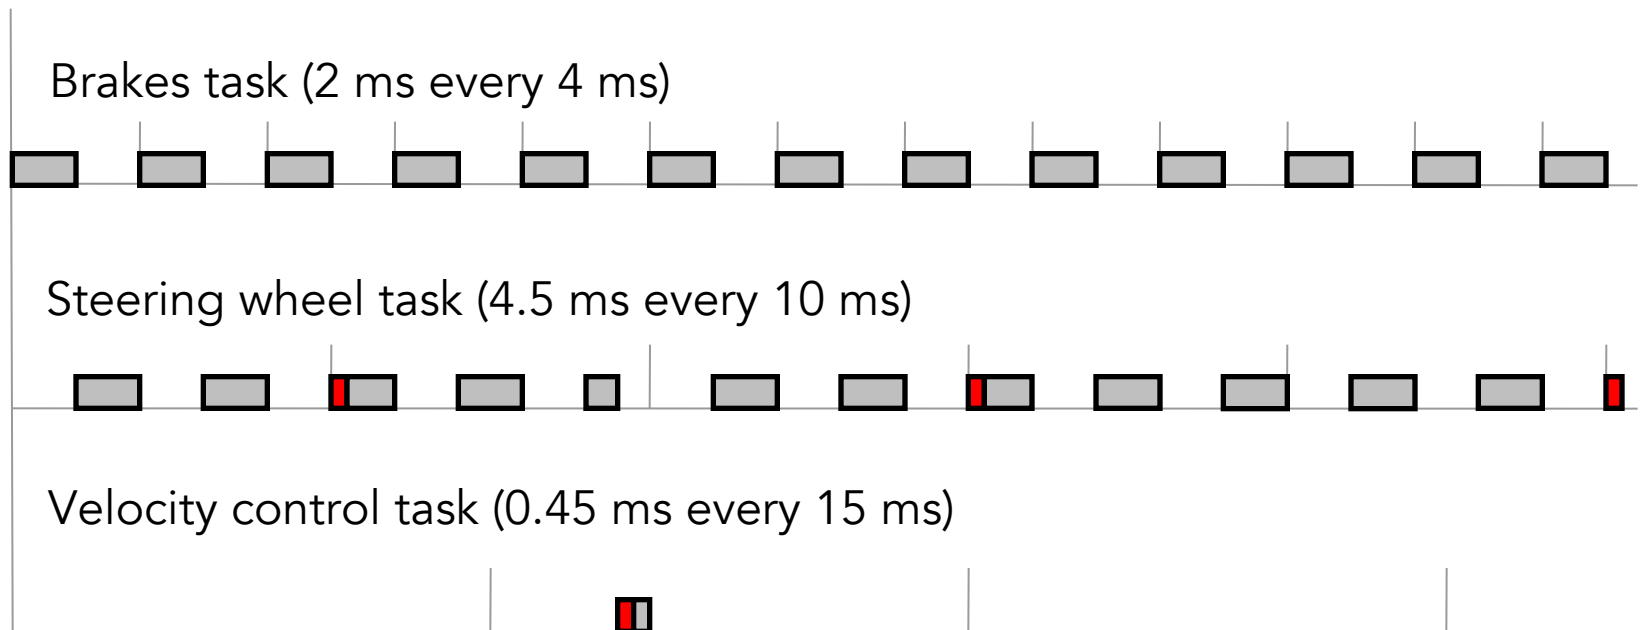
- Preemptive versus non-preemptive
  - Preemptive: Higher-priority tasks can interrupt lower-priority ones
  - Non-preemptive: They can't



**In this example, will non-preemptive scheduling work?**  
**Hint: Compare relative deadlines of tasks to execution times of others**

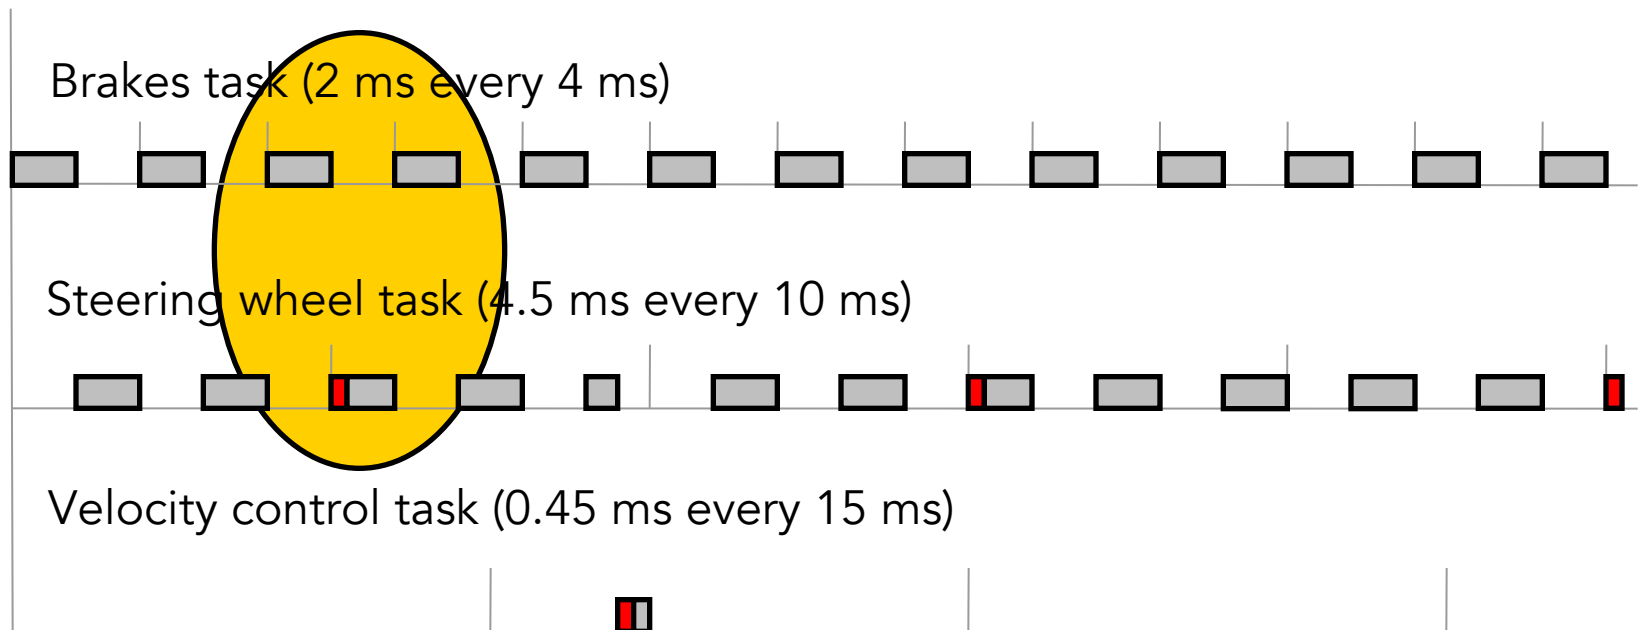
# Timeline

- Even with preemption, deadlines are missed!
- Average utilization < 100%



# Timeline


- Deadlines are missed!
- Average utilization < 100%



# Timeline

- Deadlines are missed!
- Average utilization < 100%


**Fix:**  
**Give this task invocation**  
**a lower priority**



# Timeline

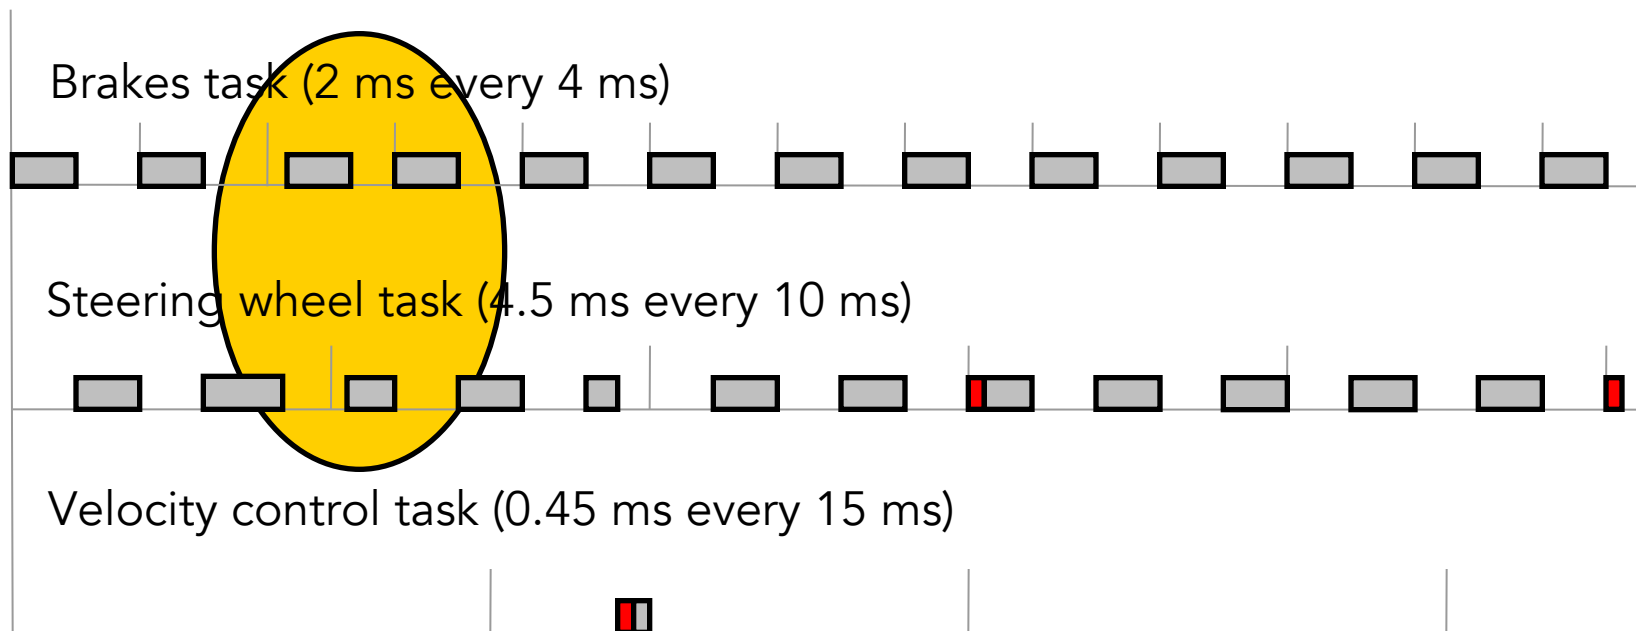
- Deadlines are missed!
- Average utilization < 100%

**Fix:**  
**Give this task invocation**  
**a lower priority**



# Task scheduling

- Static versus Dynamic priorities?
  - Static: All jobs (instances) of the same task have the same priority
  - Dynamic: Jobs (instances) of same task may have different priorities



**Intuition: Dynamic priorities offer the designer more flexibility and hence are more capable of meeting deadlines**

# Examples of policies

---

- Static priority policies
  - **Rate monotonic priority:** tasks with shorter periods get higher priority
  - **Deadline monotonic priority:** tasks with shorter deadlines get higher priority
  - Rate monotonic priorities and deadline monotonic priorities are identical if relative deadlines are equal to the periods
  - **Shortest job first policy**
- Dynamic priority policies
  - **Earliest deadline first:** jobs with the earliest absolute deadline take highest priority
  - **First In, First Out:** jobs with earliest arrival time take highest priority



# Interesting questions

---

- What is the optimal dynamic priority scheduling policy?  
(Optimal: meets all deadlines as long as any other policy in its class can)
  - Can it meet all deadlines as long as the processor is not over-utilized? [ $U \leq 1$ ]
- What is the optimal static priority scheduling policy?
  - When can it meet all deadlines?
  - Can it meet all deadline as long as the processor is not over-utilized?

# Interesting questions

---

- What is the optimal dynamic priority scheduling policy?  
(Optimal: meets all deadlines as long as any other policy in its class can)
  - Can it meet all deadlines as long as the processor is not over-utilized? [ $U \leq 1$ ]
- What is the optimal static priority scheduling policy?
  - When can it meet all deadlines?
  - Can it meet all deadline as long as the processor is not over-utilized?



**Utilization  
Bounds**

# Utilization bounds for schedulability


---

- $U_S$  is called a **utilization bound** for a given scheduling policy  $S$  if and only if
  1. All task sets with utilization  $\leq U_S$  can be scheduled using the policy  $S$  and
  2. *There exists at least one* task set with utilization  $(U + \epsilon)$  for some  $\epsilon > 0$  that **cannot** be scheduled using policy  $S$ .
- Of course, the maximum value that  $U_S$  can attain for any  $S$  is 1. Why? In class
- $U_S$  is a function of task utilizations  $u_1, \dots, u_n$

# Optimality result for EDF scheduling

---


- EDF is the optimal dynamic priority scheduling policy
  - Priorities correspond to absolute deadlines
  - It can meet all deadlines whenever the processor utilization is less than 100%
  - Intuition:
    - You have HW1 due tomorrow and HW2 due the day after, which one do you do first?
    - If you started with HW2 and met both deadlines you could have started with HW1 (in EDF order) and still met both deadlines
    - EDF can meet deadlines whenever anyone else can



# Optimality result for EDF scheduling

---

- EDF is the optimal dynamic priority scheduling policy
  - It can meet all deadlines whenever the processor utilization is less than 100%
  - Intuition:
    - You have HW1 due tomorrow and HW2 due the day after, which one do you do first?
    - If you started with HW2 and met both deadlines you could have started with HW1 (in EDF order) and still met both deadlines
    - EDF can meet deadlines whenever anyone else can



# Why study static-priority policies?

---

- EDF is the optimal dynamic scheduling policy and has a utilization bound of 1.
  - The utilization bound is 1 (or 100%) when tasks have periods equal to their relative deadlines.
- EDF, however, is hard to implement in most systems.
  - Complexity is high.
  - Job queues need to be reordered often (high overhead!).
  - Most hardware subsystems allow only static priorities.

# What you should know

---

- Definitions
  - Tasks
  - Task invocations
  - Release/arrival time,
  - Absolute deadline and relative deadline
  - Period
  - Start time and finish time
- Preemptive versus non-preemptive scheduling
- Priority-based scheduling
- Static versus dynamic priorities
- Utilization and Schedulability
- Optimality of EDF

## Another example

---

- You are the system administrator of Ameritrade.com (online stock trading site)
- You offer the following guarantee to your premium customers:
  - Stock trades of less than a \$100,000 value go through in 8 seconds or you charge no commission.
  - Stock trades of more than a \$100,000 value go through in 3 seconds or you charge no commission.
- Non-premium customers do not enjoy these guarantees
- Your job is to ensure that the premium customers are always served within their agreed-upon maximum latencies. [What needs to be done?](#)



# For later: Aperiodic tasks


---

- Periodic tasks vs. aperiodic tasks
  - Aperiodic tasks may arrive at any time
  - Periodic tasks arrive at regular intervals [strictly  $P_i$ ]
- Sporadic tasks
  - Successive arrivals have a minimum separation distance [greater than or equal to  $P_i$ ].
- How does the lack of periodicity affect scheduling, and schedulability analysis?

## For later: Precedence constraints

---

- In the discussion thus far, we focused on tasks that have no dependencies.
- What if tasks have precedence constraints?
  - Tasks can execute only if their predecessors have finished execution.



# For later: Resource constraints

---

- In addition to the CPU, tasks may need resources
  - Memory
  - Disk
  - Shared data structures
- Types of resources
  - Space multiplexed: An example is the memory system. Different tasks may use different portions of the resource.
  - Time multiplexed: Only one task can access the resource at a time. An example is a data structure protected by a lock.
- How do resource constraints affect scheduling?